

# Architectural Design Patterns for Baseline Regulation in Embodied Systems

**Author:** Tyson Jeffreys **Date:** December 30, 2025 **Scope:** conceptual supplement (not empirical).

This section presents concrete, implementable architectural patterns that instantiate the paper's core idea: robustness in embodied systems can emerge from **baseline regulation and slow, global constraint signals**, rather than from centralized correction alone.

These examples are not presented as algorithms or proofs, but as **design patterns** that can be layered onto existing robotics and embodied AI stacks.

---

## Example 1: A “Baseline Regulator” Layer Beneath RL / MPC

### Goal

Maintain a low-correction, low-energy *stable manifold* during nominal operation so that perturbations do not cascade into instability.

### Design

Introduce a slow control loop (e.g., 1–5 Hz) that continuously estimates a **regulatory cost** using signals already present in most systems, such as:

- total actuator effort (e.g., torque<sup>2</sup>)
- jerk penalties
- contact slip rate
- state-estimation uncertainty (e.g., covariance)
- prediction error between expected IMU / foot forces and observed signals

The baseline regulator does **not** command actions or trajectories.

Instead, it **biases parameters** in the existing control stack:

- adjusts impedance / stiffness gains
- adjusts MPC cost weights (e.g., stability vs. speed)
- adjusts allowed step height and speed envelopes

### Why It Matches the Paper

Stability emerges from **distributed regulation rather than centralized command**.

The planner and learned policies remain intact; the baseline regulator functions as an under-layer that prevents overcorrection spirals.

---

## Example 2: Global Constraint Signals as Broadcast Budgets

### Concept

Replace task-local optimization with a small set of slow, shared variables that every subsystem can “feel.”

### Candidate Global Signals

- **Energy budget:** battery SoC, instantaneous power draw, remaining mission energy
- **Thermal budget:** motor and driver temperatures, predicted thermal headroom
- **Traction budget:** global slip likelihood (from fused foot force, IMU, and terrain classification)
- **Compute / latency budget:** control-loop latency, dropped frames

### How to Use Them

Broadcast these signals to all layers:

- state estimator
- gait generator
- manipulation controller
- RL policy

Each layer treats them as **boundary conditions**, not commands:

- the estimator increases smoothing when compute budget is tight
- locomotion reduces aggressive maneuvers when traction budget is low
- manipulation reduces acceleration when thermal headroom is limited

This resembles whole-body control’s “respect constraints” principle, but with a critical difference: the signals are **slow, shared, and always present**, synchronizing subsystems continuously rather than reacting locally.

---

## Example 3: Predictive Coupling for Active Sensing

### Concept

Misalignment between sensing and motion increases prediction error and corrective effort.

### Concrete Robot Scenario

A mobile robot with a camera mounted on a pan–tilt unit:

- when the camera is aimed far off the velocity vector, optical flow predictions and inertial predictions diverge
- estimator strain increases

- control corrections increase downstream

## Design

Introduce a coupling objective:

- add a cost that gently biases sensor-frame alignment toward expected motion
- allow deviation only when an explicit task requires it (e.g., search or inspection)

## Engineering Metric

Minimize the discrepancy between predicted optical flow (from the motion model) and observed flow. Feed this prediction-error term into the baseline regulator as part of the regulatory cost.

This aligns with forward-model and predictive-processing traditions, but reframed as **regulatory cost minimization**, not task optimization.

---

## Example 4: Reinforcement Learning with Global Constraints as Shared Latent State

### Concept

Make global constraints first-class citizens in learning-based systems.

### Design

Introduce a small set of slow-evolving latent variables:

- energy headroom
- thermal headroom
- traction confidence
- estimator confidence

Then:

- feed these variables into the policy so action selection respects global conditions
- add auxiliary losses that encourage the policy's internal representations to predict or align with these variables

This makes global constraints **accessible across control layers**, rather than implicit or emergent.

---

## Example 5: Distributed Biasing in a Typical ROS Stack

### Minimal Intervention Pattern

Without rewriting existing systems:

1. Add a **Constraint Signal Node** that publishes `/global_constraints`

- energy
- thermal
- traction
- latency

2. Modify existing controllers to subscribe and modulate parameters such as:

- maximum velocity
- maximum joint acceleration
- impedance stiffness
- planner horizon

3. Log a single **regulatory cost** scalar to tune for low-demand stability.

This pattern directly translates the paper's design principles into a deployable architecture with minimal disruption.

---

## Summary

Across these examples, the common move is architectural rather than algorithmic:

- introduce slow, shared constraint signals
- bias existing controllers instead of commanding them
- maintain a low-correction operating manifold during nominal conditions

Planning and learning then operate **on top of coherence**, rather than being responsible for creating it.